

PR #590: both improvements, measured against main

Left: end-to-end cost of explaining n instances against a 100-row background (heloc, RandomForest 20 × depth 8), median of 3 runs. Right: aggregate_base_attributions timed alone on frozen SII inputs, median of 9 runs. Values are unchanged in both cases — the interventional path matches a brute-force oracle to 1.7e-16 and the aggregation is bit-identical. Single thread; shap and Woodelf are untouched by the PR and act as the control (1.01× / 1.04×).

